



StoryPoints AI Internship Program

Week 1 Assignment — Data Extraction & Basic Transformation

Useful Links:

Before you begin your tasks, take some time to go through the following resources. These will give you help you connect with mentors and peers.

Mentor's Profiles

Success Coach Nilesh's LinkedIn: [Link](#)

Arun Bonam's LinkedIn: [Link](#)

Community

WhatsApp Group: [Link](#) → Stay updated, ask questions, and share learning. We're currently sharing 101 Hacks to help you grow in Life, Career, Leadership, and Business. These will also support your journey as a professional tester.

Share Your First Steps:

Before diving into the Week 1 tasks, we'd love for you to share this new beginning with your network. Create a short LinkedIn post announcing that you've started your Internship with London Success Academy.

This not only celebrates your milestone but also starts building your professional brand right from day one. 🚀

Welcome to the StoryPoints AI Data Engineering Internship!

In Week 1, you will learn how to fetch data from multiple sources, perform basic transformations, and store it in a production-ready format on Google Cloud Storage (GCS). This assignment mimics a real-world scenario where data engineers extract, transform, and manage large datasets.

♦ Prerequisites

Set up Google Cloud using the following resources:

Create a Free GCP Account: <https://cloud.google.com/free>

Set Up a Project in GCP:

<https://cloud.google.com/resource-manager/docs/creating-managing-projects>

Create a GCS Bucket: <https://cloud.google.com/storage/docs/creating-buckets>

Install Google SDK CLI: <https://cloud.google.com/sdk/docs/install>

BigQuery Introduction: <https://cloud.google.com/bigquery/docs/introduction>

Scenario

StoryPoints AI processes large-scale clickstream, transaction, and currency conversion data daily. As a data engineer, your Week 1 task is to extract data from multiple sources (CSV files and APIs), apply basic transformations, and store the cleaned datasets into a structured, partitioned GCS bucket.

Datasets & API

Dataset	Format	Rows	Description
clickstream.csv	CSV	~200k	Website clickstream data
transactions.csv	CSV	~100k	Customer purchase transactions
ExchangeRate-API	JSON (API)	Live	Fetch real-time currency conversion rates

Download files from the GitHub repository:

<https://github.com/storypointsai/dataengineering-internship>

Use the free ExchangeRate-API to fetch real-time USD-based conversion rates. You need to register for an API key.

Using the ExchangeRate-API

Step 1: Register for a Free API Key: <https://www.exchangerate-api.com/>

Step 2: API Endpoint

https://v6.exchangerate-api.com/v6/YOUR_API_KEY/latest/USD

Step 3: Example Python Code

```
import os
import requests

API_KEY = "YOUR_API_KEY" # Replace with your key
url = f"https://v6.exchangerate-api.com/v6/{API_KEY}/latest/USD"

response = requests.get(url)
data = response.json()

if response.status_code == 200 and data["result"] == "success":
    rates = data["conversion_rates"]
    print("USD to INR:", rates["INR"])
else:
    print("API Error:", data)
```

Step 4: Sample JSON Response

```
{
  "result": "success",
  "base_code": "USD",
  "time_last_update_utc": "Fri, 05 Sep 2025 00:00:01 +0000",
  "conversion_rates": {
    "EUR": 0.92,
    "INR": 83.1,
    "GBP": 0.78,
    "JPY": 145.3
  }
}
```

⚡ Store raw JSON responses into: data/raw/api_currency/YYYY-MM-DD/

Assignment Tasks

1. Task 1: Explore datasets and document schemas, null values, and duplicates in README.md.
2. Task 2: Extract data — read CSVs in chunks (e.g., 50,000 rows) and fetch currency rates via API.
3. Task 3: Basic transforms — standardize column names, convert timestamps to UTC, deduplicate, and enrich transactions with amount_in_usd using conversion rates.
4. Task 4: Load — write cleaned outputs to GCS paths partitioned by ingest_date=YYYY-MM-DD/ for each dataset.
5. Task 5: Logging & alerts — log record counts and failures; print warnings if inputs are missing or API fails.
6. Task 6: Architecture diagram — draw your pipeline (sources → ETL scripts → partitioned GCS) and include it in README.

Deliverables

- Python scripts for ingestion, transformation, and GCS loading.
- README.md with dataset understanding, approach, assumptions, and architecture diagram.
- Partitioned outputs stored in GCS.
- Proper Documentation with the results of the tasks that you achieved (ex: screenshots)
-

Submission Workflow

- Fork the starter GitHub repository (link will be provided).
- Create a PRIVATE repository in your GitHub and push your code there.
- Add StoryPoints AI reviewers as collaborators. (arun.bonam@storypointsai.com)
- Work in a feature branch (e.g., week1-task-<name>) and commit frequently.
- Create a Pull Request to your main branch and assign reviewers.
- Submit your repository link.

Pro Tips

- Use pandas chunking: pd.read_csv(file, chunksize=50000) for large CSVs.
- Implement retry with exponential backoff for API calls.
- Partition outputs by ingest_date=YYYY-MM-DD/ for efficiency.
- Use UTC timestamps to avoid timezone issues.
- Keep commits small and meaningful.